

000 GPU-ACCELERATED RETRIEVAL AND SCALABLE 001 002 MODEL SERVING 003 004

005 **Anonymous authors**

006 Paper under double-blind review

007 008 ABSTRACT 009

010 This report addresses GPU-based information retrieval, focusing on distance met-
011 rics and top-K retrieval for large-scale data (Task 1), along with a queue-and-batch
012 serving design for handling real-time inference requests (Task 2). By accelerating
013 retrieval kernels on GPUs, we significantly reduce per-query latency. In parallel,
014 we adopt systematic request queueing and batching to smooth bursty traffic, in-
015 crease throughput, and maintain stable average and tail latencies. Our experiments
016 show that the synergy between low-level GPU optimizations and high-level or-
017 chestration yields robust performance improvements for real-world ML pipelines.
018 We highlight challenges, potential solutions, and future directions for scaling these
019 methods in production environments.

020 021 1 INFORMATION RETRIEVAL ON GPU (TASK 1)

022 1.1 MOTIVATION

023 Information Retrieval (IR) systems aim to efficiently locate and deliver relevant data in response
024 to user queries through three core stages: indexing, query processing, and ranking/retrieval Baeza-
025 Yates & Ribeiro-Neto (2011). The indexing stage structures data for rapid access using techniques
026 including inverted indices for text searches Zobel & Moffat (2006), spatial indices like k-d trees for
027 multidimensional data Bentley (1975), and clustering methods such as K-Means for approximate
028 searches. Large-scale IR in machine learning systems face significant computational challenges
029 due to the need to process high-dimensional data (e.g., millions of vectors with thousands of di-
030 mensions), where traditional CPU-based approaches become infeasible because of the quadratic
031 complexity of exhaustive distance calculations.

032 Clustering enhances IR systems by organizing data into semantically coherent groups, reducing
033 computational complexity. This approach allows dynamic pruning of irrelevant clusters, valuable for
034 real-time applications like recommendation systems. Modern GPU architectures address clustering
035 bottlenecks through three key mechanisms: parallelism that enables thousands of concurrent threads
036 for distance computations, optimized memory hierarchies featuring coalesced memory access and
037 shared memory utilization, and scalable processing of high-dimensional data through batched op-
038 erations. Studies demonstrate GPU-accelerated clustering achieves 20–50x speedups over CPU
039 implementations for large-scale IR tasks, though memory management remains essential for opti-
040 mal performance Johnson et al. (2021a). In large-scale recommendation systems like those used by
041 e-commerce or streaming platforms, efficient retrieval of similar user profiles or products from mil-
042 lions of high-dimensional vectors is critical. By combining clustering with GPU-accelerated search,
043 these systems can deliver real-time suggestions with significantly reduced latency and computational
044 cost (Elkahky et al. (2015)).

045 1.2 OVERVIEW: SYSTEM ARCHITECTURE

046 The Fig 1 1 illustrates a GPU-accelerated clustering-based IR pipeline optimized for large-scale,
047 high-dimensional data. The workflow begins with L2-normalization for consistent similarity mea-
048 surement. GPU-accelerated K-Means then clusters the data through parallelized distance calcu-
049 lations and iterative centroid updates. During retrieval, queries first identify relevant clusters be-
050 fore performing either exact K-Nearest Neighbor(KNN) or Approximate Nearest Neighbor(ANN)
051 search, with all computations optimized for GPU parallelism and memory efficiency.

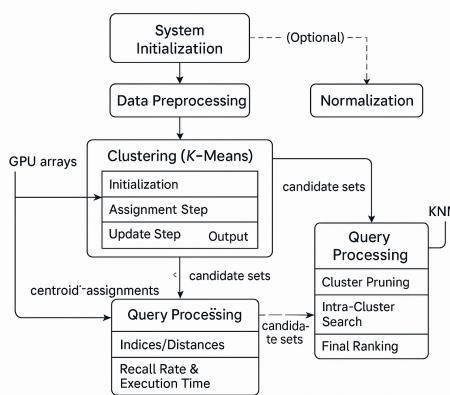


Figure 1: The system workflow.

Key Design Principles:

GPU Acceleration for Computational Efficiency CuPy for GPU-optimized operations including parallel distance calculations and centroid updates, achieving significant speedups over CPU implementations for large-scale datasets. (Johnson et al. (2021c)).

ANN-KMeans Hierarchical Retrieval: Integrates approximate nearest neighbor search with K-Means clustering by first pruning the search space through centroid comparisons, then performing exact distance calculations only within relevant clusters, optimizing the trade-off between retrieval accuracy and computational efficiency. (Zhang et al. (2019b)).

Multi-Distance Metric Support: Implements four complementary distance functions (L2, cosine, dot product, Manhattan) to handle diverse retrieval scenarios, with GPU-accelerated computations ensuring scalability across high-dimensional data. (Mikolov et al. (2013)).

1.3 DISTANCE FUNCTIONS

In large-scale IR tasks, we implemented four core distance functions—Euclidean (L2), Manhattan (L1), cosine, and dot product—using CuPy to enable efficient, GPU-accelerated vectorized computations.

Problem Definition: Distance functions quantify similarity or dissimilarity between data points in information retrieval tasks. As dataset sizes and dimensionality grow, computing these measures becomes prohibitively slow on CPU-based systems, creating a scalability bottleneck. GPU acceleration addresses this by using massive parallelism to speed up distance calculations for efficient large-scale vector similarity computations.

Design Choice: Our approach addresses this challenge by using GPU acceleration to implement four fundamental distance metrics: dot product, cosine distance, squared Euclidean (L2) distance, and Manhattan (L1) distance. The vectorized GPU operations enable efficient parallel processing and eliminating the overhead of traditional CPU-based methods. All the functions are expressed below using two vectors x and y .

Dot Product:

$$\text{Dot Product}(x, y) = \sum_{i=1}^D x_i \cdot y_i$$

The dot product quantifies the degree of alignment between two vectors. Its low computational cost makes it ideal for preliminary filtering in tasks like recommendation systems. In high-dimensional settings, it enables fast approximation of similarity before applying more refined distance metrics.

Cosine Distance:

$$\text{cosine_sim}(x, y) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$$

108 Cosine distance evaluates the angular difference between vectors by normalizing them and computing
109 the inverse of cosine similarity. It is particularly effective in scenarios like semantic search or
110 text embedding comparisons, where orientation is more informative than magnitude. This metric
111 helps mitigate the curse of dimensionality by focusing on direction rather than scale.

112 **Squared Euclidean (L2) Distance:**

114

115
$$\text{L2_distance}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^D (x_i - y_i)^2$$

116

117 L2 distance measures the geometric similarity between vectors by summing squared component-
118 wise differences. It is well-suited for clustering tasks aimed at minimizing intra-cluster variance
119 due to its sensitivity to both magnitude and direction. Its efficiency and interpretability make it a
120 common choice for image analysis and pixel-level comparisons.

121 **Manhattan (L1) Distance:**

123

124
$$\text{Manhattan Distance}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^D |x_i - y_i|$$

125

126 Manhattan distance calculates the total absolute difference across vector dimensions, capturing cu-
127 mulative dissimilarity. It is more robust to outliers than L2 distance and performs well in sparse
128 or anomaly-prone datasets. This metric is widely used in grid-based systems and fraud detection,
129 where stability across features is essential.

130 **Technical Implementation :** The distance functions mentioned above perform element-wise op-
131 erations entirely on the GPU, thereby reducing computational overhead compared to CPU-based
132 methods. For low-dimensional data ($D = 2$), our GPU implementation achieves a speedup of ap-
133 proximately 2–3×, whereas for very high-dimensional data ($D = 2^{15}$), parallelism of GPUs yielded
134 a speedup of 10–20× over CPU implementations. We addressed challenges such as managing GPU
135 memory and ensuring numerical stability during operations like vector normalization by converting
136 data to `float32` and employing efficient batch processing to maintain all computations on-device.
137 Our approach using CuPy achieves scalable and efficient distance computations.

138

139 1.4 TOP-K RETRIEVAL

140 Problem Definition: In KNN search, we find K points closest to a query point X from dataset A
141 using a distance metric. To improve efficiency, we implement ANN with clustering to partition
142 data and restrict searches to relevant clusters, effectively leveraging GPU parallelism for handling
143 millions of points. **Design Choice:** To address scalability, our design adopts a GPU-accelerated,
144 clustering-based ANN approach that first partitions data using K-Means, then performs a two-stage
145 search—selecting relevant clusters via centroids and searching within them. This strategy mini-
146 mizes distance computations and CPU-GPU transfers, enabling high throughput with tunable speed-
147 accuracy trade-offs. By narrowing search scope and leveraging parallelism, the system achieves
148 significant speedups with minimal loss in recall.

149 **Implementation:**

150 **KNN:** We ran our GPU-based KNN function and measured its execution time. As a baseline, we
151 ran an exact KNN on the CPU (using NumPy for distance computations) under the same conditions.
152 We then compared the results. Tests were repeated across different vector dimensions and dataset
153 sizes to evaluate scalability.

154 **K-Means:** The K-Means implementation is fully CuPy-based, performing all computations di-
155 rectly on the GPU for maximum efficiency. It handles edge cases like empty clusters through re-
156 initialization and returns the final results to the CPU using `.get()`. This design ensures scalability
157 while maintaining flexibility for extensions such as cosine-based clustering or performance bench-
158 marking.

159 **Approximate Nearest Neighbors (ANN): Clustering-Based ANN** is implemented, offering trade-
160 offs between speed and accuracy. In this approach, the dataset is first partitioned using GPU-
161 accelerated K-Means, forming a coarse-grained index. During a query, the input vector is compared

162 to all cluster centroids, and the top- K_1 nearest clusters are selected based on the distance metric
163 selected. All data points within these selected clusters are then gathered as candidate neighbors.
164 The system computes exact distances between the query and these candidates on the GPU. This
165 localized search reduces the computational overhead compared to a brute-force search. In addition
166 to the clustering-based approach, an additional implementation of **HNSW** (Hierarchical Navigable
167 Small World) and **IVFPQ** (Inverted File with Product Quantization) is also included.

168 **Evaluation:** Referring the table 1, At low dimensionality ($D = 2$), the GPU showed no speed ad-
169 vantage and was significantly slower due to overhead, taking approximately 11.9 seconds compared
170 to 1.06 seconds on the CPU. However, at higher dimensionality ($D = 1024$), the GPU completed
171 an L2 ANN query in about 1.72 seconds versus 1.79 seconds on the CPU. This indicates that the
172 crossover point for GPU efficiency occurs as vector dimensionality increases, with GPU perfor-
173 mance becoming competitive or superior from around $D = 1024$.

175 2 MODEL SERVING SYSTEMS (TASK 2)

177 2.1 MOTIVATION

179 Modern machine learning (ML) applications, including recommendation engines and retrieval-
180 augmented generation (RAG) systems, often handle vast numbers of inference requests under strin-
181 gent latency requirements (Hazelwood et al., 2018; Crankshaw et al., 2017). Even moderate in-
182 creases in queuing or processing delay can degrade user experiences, while poor resource man-
183 agement risks leaving GPUs underutilized or causing CPU bottlenecks (Wu et al., 2022; Ali et al.,
184 2022). These problems become especially severe when traffic arrives in bursts, making it difficult
185 to maintain consistent performance (Yu et al., 2021). Such challenges are further amplified by the
186 complexity of large models, which require substantial compute and memory resources.

187 In practice, frameworks like TensorFlow Serving (Developers, 2022), TorchServe (PyTorch Contri-
188 butors, 2025), and NVIDIA Triton (Morales-Hernández et al., 2021) employ dynamic batching and
189 request queueing to manage large-scale or bursty workloads more effectively. By gathering multi-
190 ple requests into a single forward pass on the GPU, batching reduces overhead and harnesses the
191 parallelism of modern accelerators. Meanwhile, systematic request queueing prevents momentary
192 load spikes from overwhelming the system. Together, these strategies have been shown to mitigate
193 both average and tail latencies, improving cost efficiency and scalability in production-grade ML
194 services (Yu et al., 2021). The approach explored here similarly adopts queueing and batching to
195 reduce resource waste, stabilize throughput, and meet the real-time needs of contemporary machine
learning pipelines.

197 2.2 OVERVIEW

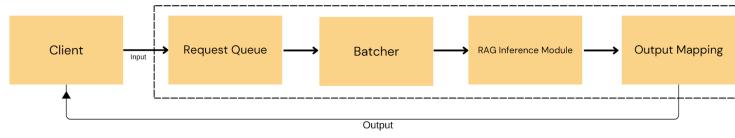
199 Our system is designed around a multi-stage workflow that integrates retrieval and large-scale infer-
200 ence, reflecting practical setups in modern ML-serving pipelines. At a high level, incoming requests,
201 often user queries, are first placed into a centralized request queue. This queue decouples request
202 arrival from processing, thereby preventing transient spikes in traffic from immediately overloading
203 the system. Once requests accumulate, a background process forms a batch of requests, aiming to
204 achieve favorable utilization by submitting them together to the GPU in a single forward pass.

205 In our implementation, each request triggers an initial retrieval step that identifies relevant docu-
206 ments or embeddings, a step often used in retrieval-augmented generation. The system then con-
207 structs a prompt or input suitable for the chosen large language model. After batched inference
208 completes, responses are split and mapped back to their respective request IDs. This design mini-
209 mizes overhead per request by reusing shared resources, including GPU kernels and memory, across
210 multiple requests simultaneously. In addition, developers can configure the maximum batch size and
211 the waiting period for building a batch, striking a balance between lower latency for small numbers
212 of incoming requests and higher throughput during bursts.

213 2.2.1 REQUEST HANDLING AND PROCESSING WORKFLOW

215 Figure 2 illustrates the overall architecture (how the process inference requests are handled), with
the core steps as follows:

- 216 1. **Request Arrival:** Clients send inference requests at varying RPS (requests per second)
 217 rates, simulating real-world traffic and load patterns.
 218
 219 2. **Request Queue:** Incoming requests are buffered in a queue to smooth out short-term
 220 spikes, avoid immediate overload and maintain order.
 221
 222 3. **Batching Module :** Requests from the queue are grouped into batches using a fixed batch
 223 size only. This allows for optimized utilization of GPU resources.
 224
 225 4. **RAG Inference Execution:** For each batch of requests, the system retrieves relevant doc-
 226 uments via their precomputed embeddings, constructs a prompt, and uses a text-generation
 227 model to produce the final answer. These batched requests are processed in a single forward
 228 pass on the GPU, leveraging hardware parallelism.
 229
 230 5. **Output Mapping:** The responses are matched with the original request IDs and returned
 231 to the respective clients.



232
 233
 234 Figure 2: Workflow Diagram: Request Queue and Batch Processing Pipeline
 235
 236

237
 238 The key advantage of this pipeline is that **queueing** helps handle bursty workloads and prevents
 239 overload, while **batching** maximizes GPU throughput by grouping multiple requests together.
 240

241 2.3 SYSTEM MEASUREMENT

242
 243 We conducted performance tests by issuing concurrent HTTP requests at controlled rates, using a
 244 Python-based load generator that varied the arrival patterns (uniform, bursty/Poisson, random) in
 245 order to emulate realistic traffic. This allowed us to observe how the system behaves under steady
 246 conditions as well as sudden spikes. To ensure comprehensive evaluation, we measured four primary
 247 metrics: **average latency**, **tail latency** (95th or 99th percentile), **throughput** (requests served per
 248 second), and **error rate**. These metrics were chosen because they reflect the system's ability to
 249 handle high loads while remaining responsive, especially when large models and bursty workloads
 250 are involved.

251
 252 Concretely, the load generator logged timestamps, response codes, and response times for every
 253 request, enabling us to compute average and percentile-based latency. We captured throughput by
 254 counting how many valid responses were returned per second, and we recorded error rates based on
 255 any failed or dropped requests. Analyzing these logs provided insight into how different configura-
 256 tions (baseline, queue-only, queue abd batcher) scaled with increasing request-per-second (RPS)
 257 loads. Our rationale for focusing on these metrics is that real-time services cannot simply maximize
 258 throughput without also guaranteeing low average and tail latencies, and error-free performance is
 259 central to system reliability.

260
 261 To collect repeatable and detailed data, the load testing script ran trials at various combinations
 262 RPS levels and load patterns for each implementation, generating summary statistics after each
 263 run. These summaries highlighted how queueing prevented the system from becoming overloaded
 264 under bursty conditions, although it introduced minor delays by buffering requests. Meanwhile,
 265 batching leveraged the GPU's parallel capacity, boosting throughput at higher loads but occasionally
 266 increasing wait times for individual requests. We faced practical challenges arising from limited
 267 hardware resources and the overhead of running both the load generator and model inference on
 268 the same machine. Consequently, we prioritized moderate RPS scenarios and shorter test windows,
 269 acknowledging that more extensive stress-testing remains desirable future work.

270 2.4 DESIGN: REQUEST QUEUE AND BATCHER
271

272 The rationale for integrating a request queue and batcher is rooted in the need to manage high vari-
273 ability in request arrival while ensuring efficient use of GPU resources. Without a queue, sudden
274 spikes can overwhelm the system, cause request failures, and leave some hardware idle under lighter
275 traffic. By buffering requests in a queue, the system can dispatch them more systematically, smooth-
276 ing out bursts. A batcher then aggregates several queued requests into a single GPU pass, reducing
277 per-request overhead such as data transfers and kernel launches. This approach balances immediate
278 responsiveness against the performance gains from larger batches; system administrators can tune
279 parameters like maximum batch size and wait time to align with specific latency or throughput goals.
280

281 Implementation involves a central queue where each incoming request is time-stamped and en-
282 queued. A background process continuously checks the queue, forming a batch when either a size
283 threshold is reached or a short timeout elapses. The system then packs all items in the batch into a
284 consolidated input for the model and runs one forward pass on the GPU. To evaluate these design
285 choices, we conducted experiments measuring throughput, latency, and error rates under different
286 traffic patterns. Our analysis showed that batching significantly boosts throughput, especially under
287 moderate to high load, although early-arriving requests in each batch can wait slightly longer.
288 Overall, queuing prevented overload by deferring arrivals rather than dropping requests, while batch-
289 ing minimized wasted GPU cycles by amortizing overhead across multiple queries. Nevertheless,
290 limited GPU memory, unpredictable arrival bursts, and the overhead of run-time scheduling posed
291 challenges. Solutions included setting a sensible batch size, implementing adaptive batch time-
292 outs, and monitoring queue lengths to detect overload conditions. Full details of these experiments,
293 along with additional metric breakdowns and extended discussion of challenges, can be found in
294 Appendix B.4.

295 2.5 OPEN DISCUSSIONS
296

297 This section reflects on the broader implications of our design and experiments, focusing on hard-
298 ware constraints, traffic variability, and potential extensions. We tested multiple arrival patterns
299 (uniform, Poisson, random) across different RPS levels, though not exhaustively, limiting definitive
300 comparisons. Each configuration was run only once, so outliers could skew results, and running
301 the load generator and inference on the same GTX 1060 (6 GB VRAM) likely inflated latencies.
302 Although queueing and batching improved throughput in moderate-load scenarios, higher RPS sat-
303 urated the GPU quickly, leaving absolute latencies higher than anticipated. In real-world appli-
304 cations, more advanced mechanisms may be necessary, particularly when traffic is highly bursty
305 or model sizes exceed GPU memory constraints. Further details, including future directions such
306 as dynamic batching, multi-threaded CPU preprocessing, and hybrid queue approaches, appear in
307 Appendix B.5.

308 3 DISCUSSION OF TASK SYNERGY
309

310 Although raw GPU acceleration (Task 1) lowers the cost of individual retrieval or distance-
311 computation kernels (Taipalus, 2024; Zhang et al., 2019a), these gains can be lost if the system
312 framework (Task 2) issues inefficient kernel launches due to poor batching or queuing. In real-
313 world pipelines, a well-optimized kernel alone does not guarantee overall throughput if frequent,
314 small-batch calls overwhelm the GPU with overhead (Crankshaw et al., 2017; Tillet et al., 2019).
315 Conversely, faster kernel routines allow Task 2 to form larger batches without causing prohibitive
316 latency. Hence, the two tasks are interdependent: Task 1 focuses on optimizing top-K searches
317 and distance metrics for efficiency (Douze et al., 2024; Malkov & Yashunin, 2020; Johnson et al.,
318 2021b), while Task 2 ensures that the parallel advantages of GPU kernels are fully realized through
319 intelligent request orchestration. Aligning these layers yields a system that combines low-level per-
320 formance with high-level scalability, necessary for robust, production-scale ML services (Li et al.,
321 2020; Zhang et al., 2019a).

324 REFERENCES
 325

- 326 Ahsan Ali, Riccardo Pincioli, Feng Yan, and Evgenia Smirni. Optimizing inference serving on
 327 serverless platforms. *Proc. VLDB Endow.*, 15(10):2071–2084, June 2022. ISSN 2150-8097. doi:
 328 10.14778/3547305.3547313. URL <https://doi.org/10.14778/3547305.3547313>.
- 329 Ricardo Baeza-Yates and Berthier Ribeiro-Neto. *Modern Information Retrieval*. Addison-
 330 Wesley, 2011. URL <https://www.pearson.com/en-us/subject-catalog/p/modern-information-retrieval/P200000000377/9780130353009>.
- 332 Jon Louis Bentley. Multidimensional binary search trees used for associative searching. *Communications of the ACM*, 18(9):509–517, 1975. URL <https://dl.acm.org/doi/10.1145/361002.361007>.
- 336 Daniel Crankshaw, Xin Wang, Giulio Zhou, Michael J. Franklin, Joseph E. Gonzalez, and Ion Stoica.
 337 Clipper: a low-latency online prediction serving system. In *Proceedings of the 14th USENIX
 338 Conference on Networked Systems Design and Implementation*, NSDI’17, pp. 613–627, USA,
 339 2017. USENIX Association. ISBN 9781931971379.
- 340 TensorFlow Developers. Tensorflow. *Zenodo*, 2022.
- 341 Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvassy, Pierre-
 342 Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. The faiss library. *arXiv
 343 preprint arXiv:2401.08281*, 2024.
- 345 Ali Mamdouh Elkahky, Yang Song, and Xiaodong He. A multi-view deep learning approach for
 346 cross domain user modeling. In *Proceedings of the 24th International Conference on World Wide
 347 Web (WWW)*, pp. 278–288. ACM, 2015. URL <https://dl.acm.org/doi/10.1145/2736277.2741667>.
- 349 Kim Hazelwood, Sarah Bird, David Brooks, Soumith Chintala, Utku Diril, Dmytro Dzhulgakov,
 350 Mohamed Fawzy, Bill Jia, Yangqing Jia, Aditya Kalro, James Law, Kevin Lee, Jason Lu, Pieter
 351 Noordhuis, Misha Smelyanskiy, Liang Xiong, and Xiaodong Wang. Applied machine learning
 352 at facebook: A datacenter infrastructure perspective. In *2018 IEEE International Symposium on
 353 High Performance Computer Architecture (HPCA)*, pp. 620–629, 2018. doi: 10.1109/HPCA.
 354 2018.00059.
- 355 Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE
 356 Transactions on Big Data*, 2021a. doi: 10.1109/TBDA.2021.3073628. URL <https://ieeexplore.ieee.org/document/9401509>.
- 359 Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE
 360 Transactions on Big Data*, 7(3):535–547, 2021b. doi: 10.1109/TBDA.2019.2921572.
- 361 Jeff Johnson et al. Gpu speedup benchmarks. *IEEE Transactions on Big Data*, 2021c. doi:
 362 10.1109/TBDA.2021.3073628. URL <https://ieeexplore.ieee.org/document/9401509>. Accessed 2024.
- 365 Yang Li, Zhenhua Han, Quanlu Zhang, Zhenhua Li, and Haisheng Tan. Automating cloud de-
 366 ployment for deep learning inference of real-time online services. In *IEEE INFOCOM 2020
 367 - IEEE Conference on Computer Communications*, pp. 1668–1677, 2020. doi: 10.1109/
 368 INFOCOM41043.2020.9155267.
- 369 Yu A. Malkov and D. A. Yashunin. Efficient and robust approximate nearest neighbor search using
 370 hierarchical navigable small world graphs. *IEEE Transactions on Pattern Analysis and Machine
 371 Intelligence*, 42(4):824–836, 2020. doi: 10.1109/TPAMI.2018.2889473.
- 372 Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word represen-
 373 tations in vector space. *arXiv preprint arXiv:1301.3781*, 2013. URL <https://arxiv.org/abs/1301.3781>.
- 376 Mario Morales-Hernández, Md B Sharif, A Kalyanapu, Sheikh K Ghafoor, Tigstu T Dullo, Suder-
 377 shan Gangrade, S-C Kao, Matthew R Norman, and Katherine J Evans. Triton: A multi-gpu open
 source 2d hydrodynamic flood model. *Environmental Modelling & Software*, 141:105034, 2021.

- 378 PyTorch Contributors. TorchServe: Serve, optimize and scale PyTorch models in production.
379 <https://github.com/pytorch/serve>, 2025. Accessed: 2025-04-13.
380
- 381 Toni Taipalus. Vector database management systems: Fundamental concepts, use-cases, and current
382 challenges. *Cognitive Systems Research*, 85:101216, 2024. ISSN 1389-0417. doi: <https://doi.org/10.1016/j.cogsys.2024.101216>. URL <https://www.sciencedirect.com/science/article/pii/S1389041724000093>.
- 383
- 384 Philippe Tillet, H. T. Kung, and David Cox. Triton: an intermediate language and compiler for
385 tiled neural network computations. In *Proceedings of the 3rd ACM SIGPLAN International
386 Workshop on Machine Learning and Programming Languages*, MAPL 2019, pp. 10–19, New
387 York, NY, USA, 2019. Association for Computing Machinery. ISBN 9781450367196. doi:
388 [10.1145/3315508.3329973](https://doi.org/10.1145/3315508.3329973). URL <https://doi.org/10.1145/3315508.3329973>.
- 389
- 390
- 391 Yongji Wu, Matthew Lentz, Danyang Zhuo, and Yao Lu. Serving and optimizing machine learning
392 workflows on heterogeneous infrastructures. *Proc. VLDB Endow.*, 16(3):406–419, November
393 2022. ISSN 2150-8097. doi: 10.14778/3570690.3570692. URL <https://doi.org/10.14778/3570690.3570692>.
- 394
- 395 Fuxun Yu, Di Wang, Longfei Shangguan, Minjia Zhang, Xulong Tang, Chenchen Liu, and Xiang
396 Chen. A survey of large-scale deep learning serving system optimization: Challenges and oppor-
397 tunities. *arXiv preprint arXiv:2111.14247*, 2021.
- 398
- 399 Chengliang Zhang, Minchen Yu, Wei Wang, and Feng Yan. MArk: Exploiting cloud services for
400 Cost-Effective, SLO-Aware machine learning inference serving. In *2019 USENIX Annual Tech-*
401 *nical Conference (USENIX ATC 19)*, pp. 1049–1062, Renton, WA, July 2019a. USENIX As-
402 sociation. ISBN 978-1-939133-03-8. URL <https://www.usenix.org/conference/atc19/presentation/zhang-chengliang>.
- 403
- 404 Xiaoxuan Zhang et al. Ann hybrid architectures. In *Proceedings of the USENIX Annual Tech-*
405 *nical Conference (ATC)*. USENIX Association, 2019b. URL <https://www.usenix.org/conference/atc19/presentation/zhang>. Accessed 2024.
- 406
- 407 Justin Zobel and Alistair Moffat. Inverted files for text search engines. *ACM Computing Sur-*
408 *veys (CSUR)*, 38(2):1–56, 2006. URL <https://dl.acm.org/doi/10.1145/1132956.1132959>.
- 409
- 410
- 411

412 A APPENDIX: TASK 1

413

414 A.1 SYSTEM WORKFLOW

415

416 The system employs a multi-stage pipeline for similarity search, encompassing data preprocessing,
417 clustering (K-Means), query processing, and evaluation. It's designed to handle high-dimensional
418 data and is optimized for execution on NVIDIA GPUs using the CuPy library. The workflow in-
419 cludes optional normalization, different approximate nearest neighbor (ANN) search techniques,
420 and exact K-Nearest Neighbors (KNN) search.

421

422 A.2 SYSTEM INITIALIZATION AND DATA PREPROCESSING

423

424 A.2.1 SYSTEM INITIALIZATION

425 The first stage, *System Initialization*, configures the CuPy environment and pre-allocates GPU
426 memory. This step is crucial for avoiding runtime memory allocation overhead. The function esti-
427 mate_max_batch_size determines the optimal batch size to prevent out-of-memory errors, given the
428 dimensionality of the input vectors and available GPU memory.

429

```
430 def estimate_max_batch_size(D, memory_limit_bytes=2 * 1024**3):  
431     bytes_per_vector = D * 4 # float32  
432     return max(1, memory_limit_bytes // bytes_per_vector)
```

432 Here, D is the dimensionality of the data, and the memory limit is set to 2GB by default. The function
433 computes how many vectors of dimension D can fit into the available GPU memory.
434

435 In the *Data Preprocessing* stage, input data (vectors A and X) are converted into CuPy arrays
436 using cp.asarray(data, dtype=cp.float32). This transfers the data to GPU memory.
437 The dimensionality of the input is also checked:

```
438 if X_cp.ndim == 1:  
439     X_cp = X_cp[None, :]  
440
```

441 This code ensures that the input vector X is always a 2D array, which is important for consistent ma-
442 trix operations later on. Additionally, the diagram shows an optional *Normalization* step, typically
443 L2 normalization, which is required when using cosine similarity:

444 Clustering (K-Means) 445

446 The *Clustering* stage employs the K-Means algorithm to group similar data points into clusters.
447 This stage involves three key steps: *Initialization*, **Assignment Step, and **Update Step*.

448 A.2.2 INITIALIZATION 449

450 Initial centroids are randomly selected from the input data. Randomly initializing K centroids from
451 the dataset ensures a diverse starting point for the algorithm. Each data point is assigned to the nearest
452 centroid based on a chosen distance metric. This step computes the distance between each data point
453 and all centroids and assigns the point to the cluster with the nearest centroid.
454

455 Centroids are updated by computing the mean of all data points assigned to each cluster. The func-
456 tion also includes a fallback mechanism for empty clusters to prevent errors. If a cluster is empty
457 (no data points assigned to it), the centroid remains unchanged to avoid introducing NaN values.
458

459 The algorithm iterates between the assignment and update steps until a convergence criterion is
460 met. The *KNN* path finds the K nearest neighbors to a query vector by exhaustively computing
461 distances between the query vector and all data points. The cp.argsort function efficiently
462 identifies the indices of the K smallest distances, which are then returned as the nearest neighbors.

463 A.2.3 APPROXIMATE SEARCH (ANN)

464 The *Approximate Search* path provides faster, albeit less precise, search results using tech-
465 niques like cluster pruning and graph-based search. The code presents three ANN search methods:
466 Cluster-based ANN (our_ann), HNSW-based ANN (our_ann_hnsw), and IVFPQ-based ANN
467 (our_ann_ivfpq).

468 **Cluster-Based ANN** This method first identifies the nearest clusters to the query and then searches
469 only within those clusters. It uses K1 to specify the number of clusters to consider and K2 to specify
470 how many points to retrieve from each cluster. Finally, it performs a final ranking to select the top
471 K results.
472

473 **HNSW-Based ANN** The HNSW (Hierarchical Navigable Small World) method constructs a
474 multi-layered graph and uses it to efficiently find the nearest neighbors.
475

476 The construction involves building a graph where each node is connected to its M nearest neighbors.
477 The search process traverses this graph, maintaining a candidate list of the best ef candidates.
478

479 A.3 EVALUATION 480

481 The final stage of the workflow is *Evaluation*. The primary metrics are **Recall Rate & Execu-
482 tion Time*. Recall rate measures the accuracy of the ANN methods compared to the exact KNN
483 results. This function calculates the ratio of relevant items retrieved to the total number of relevant
484 items. Execution time measures the indexing and query performance of the system. This function
485 benchmarks the execution time of different ANN methods, allowing a comparison of their perfor-
486 mance characteristics.

486

487

Table 1: Recall across 10 runs for different distance functions

488

489

Distance Function	Run 1	Run 2	Run 3	Run 4	Run 5	Run 6	Run 7	Run 8	Run 9	Run 10	Mean Recall
L2 (Euclidean)	0.8	0.9	0.6	0.5	0.6	0.7	0.7	0.9	0.6	0.9	0.72
Cosine	0.5	0.6	0.5	0.6	0.4	0.6	0.4	0.8	0.3	0.5	0.52
Dot Product	0.2	0.6	0.6	0.8	0.3	0.6	0.5	0.5	0.3	0.8	0.52
Manhattan (L1)	0.6	0.3	0.5	0.8	0.5	0.6	0.5	0.4	0.6	0.6	0.54

490

491

492

493

494

495

A.4 TABLE

496

497

498

499

500

Table 2: Execution Times and Speedup Ratios (CPU vs GPU) Across Methods, Distances, and Dimensions

Method	Distance	Dimension	CPU Time (s)	GPU Time (s)	Speedup Ratio
KNN	L2	D=2	0.0005	0.0010	0.52
		D=1024	0.0328	0.0093	3.53
	Cosine	D=2	0.0003	0.0012	0.27
		D=1024	0.0051	0.0031	1.65
	Dot	D=2	0.0002	0.0009	0.26
		D=1024	0.0046	0.0025	1.81
	Manhattan	D=2	0.0002	0.0009	0.21
		D=1024	0.0072	0.0026	2.79
	ANN	L2	0.5356	11.8861	0.05
		D=1024	0.7502	2.4130	0.31
		Cosine	0.7384	12.3839	0.07
			0.4556	1.2920	0.35
		Dot	1.8328	31.3044	0.06
			0.2447	0.6975	0.35
		Manhattan	0.4970	7.7278	0.05
			0.4427	1.0271	0.43

523

524

525

526

B APPENDIX: TASK 2

527

528

529

B.1 MOTIVATION

530

531

532

533

534

535

536

537

538

539

Queueing and batching are not merely tactical optimizations; they serve as pillars for building robust, high-throughput model-serving pipelines. Recent studies illustrate how managed queues buffer incoming workloads to avoid sudden overload, thereby smoothing out utilization over time (Yu et al., 2021). Batching further boosts efficiency when infrequent or sporadic requests would otherwise leave GPUs idle. Systems like TensorFlow Serving (Developers, 2022) and TorchServe (PyTorch Contributors, 2025) incorporate these ideas by default, demonstrating real-world viability. NVIDIA Triton extends this paradigm with advanced multi-model support and adaptive scheduling (Morales-Hernández et al., 2021). The ability to handle abrupt surges or lulls in traffic while keeping large models fully utilized distinguishes state-of-the-art serving infrastructures from less flexible solutions. When combined with optimized retrieval and postprocessing stages, queueing and batching help maintain steady, predictable performance across an entire ML inference pipeline.

540 B.2 OVERVIEW
541

542 The system architecture draws inspiration from established serving frameworks such as TensorFlow
543 Serving (Developers, 2022) and TorchServe (PyTorch Contributors, 2025), both of which employ
544 internal queues and batch processing to handle inference tasks. NVIDIA Triton (Morales-Hernández
545 et al., 2021) further generalizes this approach by supporting multiple backends and dynamic batching
546 strategies that adapt to varying traffic patterns. Our focus remains on effectively dispatching requests
547 to a single model at high rates, although the same concepts of request queueing and batch scheduling
548 could be extended to multi-model serving scenarios. Separating arrival from processing not only
549 helps prevent transient overloads but also allows for more precise control over how batches are
550 formed. The trade-off, as explored in our experiments, is that batching can introduce additional wait
551 time for early arrivals, a cost that is offset by substantially improved GPU utilization and throughput
552 when request volume is high (Yu et al., 2021).
553

554 B.3 SYSTEM MEASUREMENT
555

556 In each test, a Python script submitted requests at a target rate for a fixed duration (e.g., 60 seconds),
557 while uniformly or randomly spacing out arrivals (Yu et al., 2021). To approximate real-world
558 scenarios, we varied the RPS in increments and observed when the system approached saturation.
559 Latency distributions (including 95th and 99th percentiles) were extracted from the request logs, and
560 throughput was computed by dividing the total number of successful responses by the test duration.
561 We also tracked any unsuccessful or timed-out responses to calculate error rates, which occasionally
562 occurred when request spikes coincided with GPU-intensive operations. While repeating each test
563 multiple times would have minimized variance, constraints such as GPU time quotas and hardware
564 availability limited the number of repeated runs. Nonetheless, the recorded data confirm our central
565 findings: batching substantially improves resource utilization at higher loads, and queueing helps
566 prevent instantaneous overload, although these mechanisms can introduce minor latency overheads
567 for users who arrive just as a batch is forming (Morales-Hernández et al., 2021; Developers, 2022;
568 PyTorch Contributors, 2025).
569

570 B.4 DESIGN: REQUEST QUEUE AND BATCHER
571

572 This section expands on the five primary elements of the request queue and batcher subsystem:
573 rationale, implementation details, experiments, overall findings, and challenges with corresponding
574 solutions.
575

576 **1) Rationale:** In practice, ML serving infrastructure must handle diverse and at times unpredictable
577 traffic levels. A request queue allows the system to buffer incoming requests when the GPU is busy,
578 preventing immediate rejections. By decoupling arrival from processing, the system also gains time
579 to consider how best to form batches.
580

581 **2) Implementation Details:** The queue is maintained in main memory, and each request is stored
582 alongside its arrival timestamp. A batcher thread or process polls the queue for items. If the queue
583 exceeds a size threshold, or a specified timeout occurs, the batcher builds a consolidated tensor
584 input to feed the model. After inference, responses are paired back to the original requests and
585 returned to the user.
586

587 **3) Experiments and Metric Analysis:** We ran controlled tests simulating multiple traffic patterns
588 (e.g., uniform, bursty/Poisson, random) at increasing request-per-second rates. Metrics included
589 average and tail latencies, throughput, and error rates. Repeated trials revealed how different
590 batch sizes and timeouts influenced both performance and the responsiveness experienced by
591 early-arriving requests.
592

593 **4) Overall Experimental Findings:** Batching consistently improved GPU utilization and
594 throughput, especially at moderate to high load levels. Smaller batches reduced waiting time
595 per request but led to modest throughput gains. Larger batches or extended timeouts provided
596 stronger throughput improvements but increased tail latency. Queueing successfully prevented
597 overloads when traffic spiked, allowing the system to gradually drain bursts instead of dropping requests.
598

599 **5) Challenges and Solutions:** Key challenges arose from limited GPU memory for extremely large
600 batches, difficulties in selecting a default batch size across varying workloads, and potential over-
601 head in forming batches on the fly. We addressed these by bounding the maximum batch size,
602

Implementation	Total Test Time (s)	Total Requests	Successful Requests	Error Rate (%)	Avg. Latency (s)	Throughput (req/s)
base	3281.73	53400.0	5627.0	63.023	146.0592	0.011
queue	3282.32	53400.0	6028.0	60.905	143.1135	0.011
rq+batcher	3221.68	53400.0	7644.0	53.829	140.4009	0.011

Table 3: Summary Statistics by Implementation. Each row reports metrics over the entire test duration for one system configuration.

adopting a short default timeout for low-traffic scenarios, and continuously monitoring queue length to detect overload conditions. Future work could refine these parameters adaptively, shifting thresholds based on recent demand and measured latency, thus improving both reliability and resource usage.

B.5 EXTENDED OPEN DISCUSSIONS

Limitations in Traffic Patterns and RPS Combinations: Although multiple arrival patterns were tested (uniform, Poisson, random), not every pattern was paired with every load level. As a result, it is difficult to say with certainty which pattern creates the greatest difficulty at high RPS. Comprehensive testing is time-intensive and was constrained by our hardware resources.

Single-Run Evaluations: Each configuration was run only once, so outliers might skew results. In a more rigorous setting, one would conduct multiple runs and compute confidence intervals to draw stronger conclusions about the system’s consistency.

Resource Constraints and Measurement Artifacts: By co-locating the load generator and inference on the same machine, CPU contention likely inflated latencies. Our GTX 1060 GPU (6 GB VRAM) further restricted the batch sizes we could feasibly test. Ideally, we would separate load generation onto another host and adopt a more powerful GPU to validate the concurrency solutions at higher scales.

GPU Limitations on a GTX 1060 (6GB VRAM): Although our queue-and-batcher approach did improve throughput, the absolute gains remained modest due to memory constraints and model size. Larger batches risked out-of-memory errors, but using half-precision or INT8 inference and employing model distillation could enable higher concurrency without exhausting VRAM.

Real-World Considerations and Traffic Patterns: Bursty or Poisson traffic is common in production, and our tests showed that moderate Poisson loads saw efficiency gains from batching, but higher RPS saturated the GPU quickly. In practice, advanced scaling mechanisms might be required to manage such peaks effectively.

Expected vs. Surprising Outcomes: We anticipated batching to raise throughput, though the extent of these gains was smaller than hoped (e.g., from 0.01 req/s to 0.02 req/s in some low-load scenarios). The most surprising result was that absolute latencies stayed relatively high at moderate RPS, suggesting GPU overhead and model complexity were the dominant factors.

Statistical Validity and Further Steps: Comparisons across traffic patterns remain preliminary because we ran more Poisson tests than uniform or random. Systematic tests of each pattern at each RPS level, combined with multiple runs, would clarify these differences. Tracking standard deviations or confidence intervals would give additional perspective on performance variance.

Future Directions: Dynamic Batching could adapt the batch size in real time based on observed load. **Multi-Thread CPU Preprocessing** might accelerate tokenization or embedding steps. A **Hybrid or Tiered Queue** approach could separate low-latency from high-throughput workloads. Such refinements would allow queue-and-batcher systems to serve a broader range of ML inference demands with better latency control and resource management.